


Introduction to **Information Retrieval**

Lecture 18: Link analysis

Today's lecture

- Anchor text
- Link analysis for ranking
 - Pagerank and variants

The Web as a Directed Graph




Assumption 1: A hyperlink between pages denotes author perceived relevance (quality signal)

Assumption 2: The text in the anchor of the hyperlink describes the target page (textual context)

Anchor Text

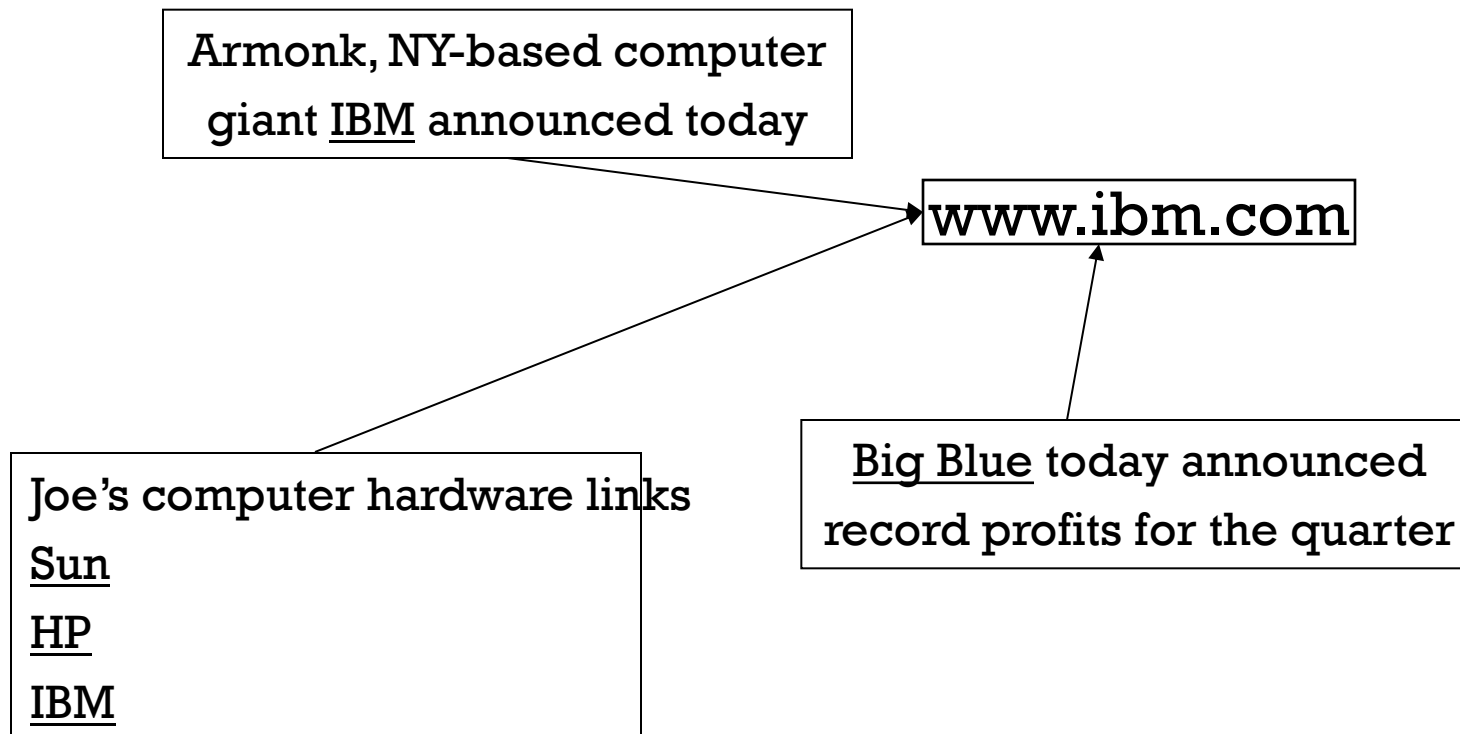
WWW Worm - McBryan [Mcbr94]

- For **ibm** how to distinguish between:
 - IBM's home page (mostly graphical)
 - IBM's copyright page (high term freq. for 'ibm')
 - Rival's spam page (arbitrarily high term freq.)



Indexing anchor text

- When indexing a document D , include anchor text from links pointing to D .



Indexing anchor text

- Can sometimes have unexpected side effects - *e.g., evil empire.*
- Can score anchor text with weight depending on the authority of the anchor page's website
 - E.g., if we were to assume that content from cnn.com or yahoo.com is authoritative, then trust the anchor text from them

Anchor Text


- Other applications
 - Weighting/filtering links in the graph
 - Generating page descriptions from anchor text

Citation Analysis

- Citation frequency
- Co-citation coupling frequency
 - Cocitations with a given author measures “impact”
 - Cocitation analysis
- Bibliographic coupling frequency
 - Articles that co-cite the same articles are related
- Citation indexing
 - Who is this author cited by? (Garfield 1972)
- Pagerank preview: Pinski and Narin '60s

Query-independent ordering

- First generation: using link counts as simple measures of popularity.
- Two basic suggestions:
 - Undirected popularity:
 - Each page gets a score = the number of in-links plus the number of out-links ($3+2=5$).
 - Directed popularity:
 - Score of a page = number of its in-links (3).



Query processing

- First retrieve all pages meeting the text query (say ***venture capital***).
- Order these by their link popularity (either variant on the previous slide).
- More nuanced – use link counts as a measure of static goodness (Lecture 7), combined with text match score

Spamming simple popularity

- *Exercise:* How do you spam each of the following heuristics so your page gets a high score?
 1. Each page gets a static score = the number of in-links plus the number of out-links.
 2. Static score of a page = number of its in-links.


Ideas of Pagerank

- Inlinks as votes
 - www.stanford.edu has 23,400 inlinks
 - www.joe-schmoe.com has 1 inlink
- Web pages are not equally “important”
 - www.joe-schmoe.com → p1
 - vs. www.stanford.edu → p2
- Are all inlinks equal?
 - Recursive question!

Pagerank scoring

- Imagine a browser doing a *random walk* on web pages:

- Start at a random page




- At each step, go out of the current page along one of the links on that page, equiprobably

- “In the steady state” each page has a long-term visit rate - use this as the page’s score.

Example – the Simple “Flow” Model

The web in 1839



$$y = y/2 + a/2$$

$$a = y/2 + m$$


$$m = a/2$$

Solving the flow equations

- 3 equations, 3 unknowns, no constants
 - No unique solution
 - All solutions equivalent modulo scale factor
- Additional constraint forces uniqueness
 - $y+a+m = 1$
 - $y = 2/5, a = 2/5, m = 1/5$
- Gaussian elimination method works for small examples, but we need a better method for large graphs

Example – the Simple “Flow” Model

The web in 1839



$$y_{new} = y_{old} / 2 + a_{old} / 2$$

$$a_{new} = y_{old} / 2 + m_{old}$$

$$m_{new} = a_{old} / 2$$

$$y = 1/3 \quad 1/3 \quad 5/12 \quad \dots \quad 2/5$$


$$a = 1/3 \quad 1/2 \quad 1/3 \quad \dots \quad 2/5$$

$$m = 1/3 \quad 1/6 \quad 1/4 \quad \dots \quad 1/5$$

Matrix-based characterization of the computation is simpler and more useful for the general case.


Markov chains

- A Markov chain consists of n states, plus an $n \times n$ transition probability matrix \mathbf{P} .
- **At each step, we are in exactly one of the states.**
- For $1 \leq i, j \leq n$, the matrix entry P_{ij} tells us the probability of j being the next state, given we are currently in state i .



Markov chains


- Clearly, for all i , $\sum_{j=1}^n P_{ij} = 1$.
- Markov chains are abstractions of random walks.



	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	1	0

Ergodic Markov chains

- A Markov chain is ergodic if
 - you have a path from any state to any other
 - For any start state, after a finite transient time T_0 , the probability of being in any state at a fixed time $T > T_0$ is nonzero.



Ergodic Markov chains

- For any ergodic Markov chain, there is a unique long-term visit rate for each state.
 - *Steady-state probability distribution.*
- Over a long time-period, we visit each state in proportion to this rate.
- It doesn't matter where we start.

Probability vectors

- A probability (row) vector $\mathbf{x} = (x_1, \dots, x_n)$ tells us where the walk is at any point.
- E.g., $(\underset{1}{000}\dots\underset{i}{1}\dots\underset{n}{000})$ means we're in state i .

More generally, the vector $\mathbf{x} = (x_1, \dots, x_n)$ means the walk is in state i with probability x_i .

$$\sum_{i=1}^n x_i = 1.$$

Change in probability vector

- If the probability vector is $\mathbf{x} = (x_1, \dots, x_n)$ at this step, what is it at the next step?
- Recall that row i of the transition prob. Matrix \mathbf{P} tells us where we go next from state i .
- So from \mathbf{x} , our next state is distributed as \mathbf{xP} .


How do we compute this vector?

- Let $\mathbf{a} = (a_1, \dots, a_n)$ denote the row vector of steady-state probabilities.
- If our current position is described by \mathbf{a} , then the next step is distributed as \mathbf{aP} .
- But \mathbf{a} is the steady state, so $\mathbf{a} = \mathbf{aP}$.
- Solving this matrix equation gives us \mathbf{a} .
 - So \mathbf{a} is the (left) eigenvector for \mathbf{P} .
 - (Corresponds to the “principal” eigenvector of \mathbf{P} with the largest eigenvalue.)
 - Transition probability matrices always have largest eigenvalue 1.

One way of computing \mathbf{a}

- Recall, regardless of where we start, we eventually reach the steady state \mathbf{a} .
- Start with any distribution (say $\mathbf{x}=(1/n, 1/n, \dots, 1/n)$).
- After one step, we're at \mathbf{xP} ;
- after two steps at \mathbf{xP}^2 , then \mathbf{xP}^3 and so on.
- “Eventually” means for “large” k , $\mathbf{xP}^k = \mathbf{a}$.
- Algorithm: multiply \mathbf{x} by increasing powers of \mathbf{P} until the product looks stable.

Power Iteration Example



P

	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	1	0

$$y_{new} = y_{old} / 2 + a_{old} / 2$$

$$a_{new} = y_{old} / 2 + m_{old}$$


$$m_{new} = a_{old} / 2$$

	\mathbf{x}_0	\mathbf{x}_1	\mathbf{x}_2	\mathbf{x}_3		\mathbf{x}_t
y	1/3	1/3	5/12	3/8		2/5
a	1/3	1/2	1/3	11/24	...	2/5
m	1/3	1/6	1/4	1/6		1/5

Spider traps

- A group of pages is a **spider trap** if there are no links from within the group to outside the group
 - Random surfer gets trapped
- Spider traps violate the conditions needed for the random walk theorem

Microsoft becomes a spider trap




	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	0	1

y		1/3	1/3	1/4	5/24		0
a	=	1/3	1/6	1/6	1/8	...	0
m		1/3	1/2	7/12	2/3		1

Random teleports

- The Google solution for spider traps
- At each time step, the random surfer has two options:
 - With probability β , follow a link at random
 - With probability $1-\beta$, jump to some page uniformly at random
 - Common values for β are in the range 0.8 to 0.9
- Surfer will teleport out of spider trap within a few time steps


Random teleports ($\beta = 0.8$)



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1/2 \\ 0 & 0 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 7/15 \\ 1/15 & 1/15 & 13/15 \end{bmatrix}$$

Random teleports ($\beta = 0.8$)



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix}$$

$$+ 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$


$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 7/15 \\ 1/15 & 1/15 & 13/15 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} = \begin{matrix} 0.33 & 0.33 & 0.28 & 0.26 & & 7/33 \\ 0.33 & 0.20 & 0.20 & 0.18 & \dots & 5/33 \\ 0.33 & 0.47 & 0.52 & 0.56 & & 7/11 \end{matrix}$$

Dead ends

- Pages with no outlinks are “dead ends” for the random surfer
 - Nowhere to go on next step
- Especially common for Web Search Engines
 - URLs that have not yet been crawled

Microsoft becomes a dead end



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1/2 \\ 0 & 0 & 0 \end{bmatrix}$$

$$+ 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 7/15 \\ 1/15 & 1/15 & 1/15 \end{bmatrix}$$

Non-stochastic!

$$\begin{matrix} y \\ a \\ m \end{matrix} = \begin{bmatrix} 0.33 & 0.33 & 0.262 & 0.216 & & 0 \\ 0.33 & 0.20 & 0.182 & 0.143 & \dots & 0 \\ 0.33 & 0.20 & 0.129 & 0.111 & & 0 \end{bmatrix}$$

Dealing with dead-ends

- Teleport
 - Follow random teleport links **with probability 1.0** from dead-ends
 - Adjust matrix **accordingly**
 - **How?**
- (Suggested by Google) prune and propagate
 - Preprocess the graph to eliminate dead-ends
 - Might require multiple passes
 - Compute page rank on reduced graph
 - **Approximate** values for deadends by propagating values from reduced graph

Q: Why approximate values and why errors are insignificant?

Pagerank summary

- Preprocessing:

- Given graph of links, build matrix \mathbf{P} .
- From it compute \mathbf{a} .
 - \mathbf{a} is the principle eigen vector of a matrix $\tilde{\mathbf{P}}$

$$\tilde{\mathbf{P}} = (1 - \beta)\mathbf{P} + \beta\mathbf{T}, \quad \mathbf{T}_{i,j} = \frac{1}{n}$$

- The entry a_i is a number between 0 and 1: the pagerank of page i .

- Query processing:

- Retrieve pages meeting query.
- Rank them by their pagerank.
- Order is *query-independent*.

Resources

- IIR Chap 21
- <http://www2004.org/proceedings/docs/1p309.pdf>
- <http://www2004.org/proceedings/docs/1p595.pdf>
- <http://www2003.org/cdrom/papers/refereed/p270/kamvar-270-xhtml/index.html>
- <http://www2003.org/cdrom/papers/refereed/p641/xhtml/p641-mccurley.html>